

The Future of Tech Training has Arrived



NEW



Hybrid Training

SELF-PACED COURSES + **LIVE** Q&A DISCUSSIONS

LEARN MORE @ ardanlabs.com/hybrid



Ultimate Rust:
Foundations

OCTOBER 2023



Use our special code for conference attendees
& receive **\$100 OFF!**
Code: **HYBRIDGC23**



Rust Web Service Deep Dive

CRUD with SQLX + Axum

herbert.wolverson@ardanlabs.com



About Herbert Wolverson

- Ardan Labs Rust Trainer & Consultant
- Author of *Hands-on Rust* and *Rust Brain Teasers*
- Author of the *Rust Roguelike Tutorial*
- Lead developer, *LibreQoS*, *bracket-lib*.
- Contributor to many open source projects.



The Pragmatic Programmers

Hands-on Rust

Effective Learning through
2D Game Development and Play



Web Service Layers

We're going to build a full stack in an hour!



Database Layer

SQLite (in memory)




Data Model

SQLx



REST API

Axum



Web View

JavaScript + HTML



Deploy

Docker





Part 1: The Database (SQLite)





Project Setup

1. Create a new Rust project with `cargo init webinar_axum`
2. Install SQLX-CLI with `cargo install sqlx-cli`
3. Create a file named `.env`
 - a. It contains `DATABASE_URL="sqlite::memory:"`
4. Create a migration with `sqlx migration add initial`

This sets up an in-memory database, and creates a file named `(timestamp)initial.sql` - ready for your database schema.





Populate the Migration File

```
CREATE TABLE books (  
    id INTEGER PRIMARY KEY AUTOINCREMENT,  
    title TEXT,  
    author TEXT  
);  
  
INSERT INTO books (title, author) VALUES ('Hands-on Rust', 'Wolverson, Herbert');  
INSERT INTO books (title, author) VALUES ('Rust Brain Teasers', 'Wolverson, Herbert');
```





Part 2: The Data Model (SQLX)



Add Dependencies



Add dependencies to Cargo.toml:

```
[package]
name = "webinar_axumcrud"
version = "0.1.0"
edition = "2021"

[dependencies]
tokio = { version = "1.32.0", features = ["full"] }
anyhow = "1.0.75"
dotenv = "0.15.0"
serde = { version = "1.0.188", features = ["derive"] }
sqlx = { version = "0.7.2", features = ["runtime-tokio", "sqlite"] }
axum = "0.6.20"

[dev-dependencies]
axum-test-helper = "0.3.0"
```

Tokio provides the async runtime.

Anyhow for simple error handling.

Dotenv for reading ".env" files.

Serde for Serialization/Deserialization.

SQLX for the Database.

Axum for the REST API (we'll need it later).

Axum-test-helper for convenient unit tests.





Create a skeleton main.rs

Setup an async Tokio main function:

```
mod db;
use crate::db::init_db;
use anyhow::Result;

#[tokio::main]
async fn main() -> Result<()> {
    // Load environment variables from .env if available
    dotenv::dotenv().ok();

    // Initialize the database and obtain a connection pool
    let connection_pool = init_db().await?;

    Ok(())
}
```





Create the db.rs module file and Book Structure

Create a new file, `db.rs` in the `src` directory.

This will contain the database API.

Let's start by defining a `Book` structure, representing a book from the database.

Deriving `Serialize` and `Deserialize` connects the structure to `Serde` — allowing for seamless transformation to/from JSON and other formats.

`FromRow` is an `SQLX` helper that makes it easy to `SELECT` data into a structure.

```
#[derive(Debug, Serialize, Deserialize, FromRow)]
pub struct Book {
    /// The book's primary key ID
    pub id: i32,
    /// The book's title
    pub title: String,
    /// The book's author (surname, lastname - not enforced)
    pub author: String,
}
```





Create the connection pool

Provides initialization for both tests and the main program.

Retrieves the db URL from environment.

Creates a connection pool.

Runs all SQLX migrations.

Since we're using in-memory - it will run each time.

```
pub async fn init_db() -> Result<SqlitePool> {  
    let database_url : String = std::env::var( key: "DATABASE_URL"?);  
    let connection_pool : Pool<Sqlite> = SqlitePool::connect(&database_url).await?;  
    sqlx::migrate!().run( migrator: &connection_pool).await?;  
    Ok(connection_pool)  
}
```






Building your Database API: Reading Data

SQLX's FromRow derivation allows you to use `query_as` to select data straight into a `Book` structure.

Data-binding prevents SQL Injection attacks.

```
pub async fn all_books(connection_pool: &SqlitePool) -> Result<Vec<Book>> {  
    Ok(  
        sqlx::query_as::<_, Book>{ sql: "SELECT * FROM books ORDER BY title,author"  
        }.fetch_all(executor: connection_pool)  
        .await?,  
    )  
}
```

```
pub async fn book_by_id(connection_pool: &SqlitePool, id: i32) -> Result<Book> {  
    Ok(sqlx::query_as::<_, Book>{ sql: "SELECT * FROM books WHERE id=$1"  
    }.bind(value: id)  
    .fetch_one(executor: connection_pool)  
    .await?)  
}
```





Creating Records

SQLX isn't an ORM (Object Relational Mapper)
- so it doesn't map structures back to the database.

Once again, binding makes the query safe and easy.

We use `get(0)` at the end to retrieve the new primary key we return from the query.

```
pub async fn add_book<S: ToString>(  
    connection_pool: &SqlitePool,  
    title: S,  
    author: S,  
) -> Result<i32> {  
    let title : String = title.to_string();  
    let author : String = author.to_string();  
    Ok(  
        sqlx::query( sql: "INSERT INTO books (title, author) VALUES ($1, $2) RETURNING id")  
            .bind( value: title)  
            .bind( value: author)  
            .fetch_one( executor: connection_pool)  
            .await?  
            .get( index: 0),  
    )  
}
```





Updating & Deleting Records

Once again, we use data binding to keep the queries safe and easy to create.

We now have a full CRUD API: Create, Read, Update, Delete.

```
pub async fn update_book(connection_pool: &SqlitePool, book: &Book) -> Result<()> {  
    sqlx::query( sql: "UPDATE books SET title=$1, author=$2 WHERE id=$3")  
        .bind( value: &book.title)  
        .bind( value: &book.author)  
        .bind( value: &book.id)  
        .execute( executor: connection_pool)  
        .await?;  
    Ok(())  
}
```

```
pub async fn delete_book(connection_pool: &SqlitePool, id: i32) -> Result<()> {  
    sqlx::query( sql: "DELETE FROM books WHERE id=$1")  
        .bind( value: id)  
        .execute( executor: connection_pool)  
        .await?;  
    Ok(())  
}
```





Unit Test the Data Model

Since we're using an in-memory database, we get the luxury of a fresh database every time.

SQLX includes a number of helpers for unit testing: fixtures, migrations.

In this case, we're keeping it simple.

(Dive into the source code to see the other tests)

Good news: All the unit tests succeeded!

```
#[cfg(test)]
mod test {
    use super::*;

    Herbert Wolverson
    #[sqlx::test]
    async fn get_all() {
        dotenv::dotenv().ok();
        let cnn : SqlitePool = init_db().await.unwrap();
        let all_rows : Vec<Book> = all_books( connection_pool: &cnn).await.unwrap();
        assert!(!all_rows.is_empty());
    }
}
```





Part 3: REST API (Axum)





Setup the REST API

Create a new file `rest.rs` in the `src` directory.

Add `mod rest` to `main.rs` to include it.

Start by defining a function that exposes our REST API.

Routers are Axum's way of connecting URLs with a handler and a function.

- A get handler handles GET requests. `post`, `patch` and `delete` handle the associated HTTP verb of the same name.
- A handler takes a function as a parameter - and runs that function when a request matches.



```
pub fn books_service() -> Router {  
    Router::new()  
        .route( path: "/", method_router: get( handler: get_all_books)) : Router<(), Body>  
        .route( path: "/:id", method_router: get( handler: get_book)) : Router<(), Body>  
        .route( path: "/add", method_router: post( handler: add_book)) : Router<(), Body>  
        .route( path: "/edit", method_router: put( handler: update_book)) : Router<(), Body>  
        .route( path: "/delete/:id", method_router: delete( handler: delete_book))  
}
```



Connect the REST API to main

Create a router function that builds another Router. Using `nest_service`, we can attach another service with a base URL.

Calling `.layer` allows us to insert the connection pool as a “Tower Layer” into Axum - available for dependency injection into handlers.

So we:

1. Connect our database API.
2. Build a router with a connection pool.
3. Listen on port 3001
4. Start the web server.



```
fn router(connection_pool: SqlitePool) -> Router {  
  Router::new()  
    // Nest service allows you to attach another router to a URL base.  
    // "/" inside the service will be "/books" to the outside world.  
    .nest_service( path: "/books", service: rest::books_service()) :Router<(), Body>  
    // Add the connection pool as a "layer", available for dependency injection.  
    .layer(Extension(connection_pool))  
}
```

```
#[tokio::main]  
async fn main() -> Result<()> {  
  // Load environment variables from .env if available  
  dotenv::dotenv().ok();  
  
  // Initialize the database and obtain a connection pool  
  let connection_pool : SqlitePool = init_db().await?;  
  
  // Initialize the Axum routing service  
  let app : Router = router(connection_pool);  
  
  // Define the address to listen on (everything)  
  let addr : SocketAddr = SocketAddr::from( pieces: ([0, 0, 0, 0], 3001));  
  
  // Start the server  
  axum::Server::bind(&addr)  
    .serve( make_service: app.into_make_service())  
    .await?;  
  
  Ok(())  
}
```



Reading Books via REST

Get_all_books and get_book are thin wrappers over our database API.

Note the `Extension(cnn) :`
`Extension<SqlitePool>` line. The handler will automatically have the connection pool injected when called.

Likewise, `Path(id)` will automatically extract from the named path entry in the route.



```
async fn get_all_books(
    Extension(cnn): Extension<SqlitePool>,
) -> Result<Json<Vec<Book>>, StatusCode> {
    if let OK(books : Vec<Book>) = all_books( connection_pool: &cnn).await {
        OK(Json(books))
    } else {
        Err(StatusCode::SERVICE_UNAVAILABLE)
    }
}
```

```
async fn get_book(
    Extension(cnn): Extension<SqlitePool>,
    Path(id): Path<i32>,
) -> Result<Json<Book>, StatusCode> {
    if let OK(book : Book) = book_by_id( connection_pool: &cnn, id).await {
        OK(Json(book))
    } else {
        Err(StatusCode::SERVICE_UNAVAILABLE)
    }
}
```

Creating, Updating and Deleting Books

Once again, we just make a thin wrapper around the database model.

Congratulations, you've implemented Create, Read, Update and Delete REST functions.

```
async fn add_book(
  Extension(cnn): Extension<SqlitePool>,
  extract::Json(book): extract::Json<Book>,
) -> Result<Json<i32>, StatusCode> {
  if let Ok(new_id :i32) = crate::db::add_book( connection_pool: &cnn, &book.title, &book.author).await {
    Ok(Json(new_id))
  } else {
    Err(StatusCode::SERVICE_UNAVAILABLE)
  }
}
```

```
async fn update_book(
  Extension(cnn): Extension<SqlitePool>,
  extract::Json(book): extract::Json<Book>,
) -> StatusCode {
  if crate::db::update_book( connection_pool: &cnn, &book).await.is_ok() {
    StatusCode::OK
  } else {
    StatusCode::SERVICE_UNAVAILABLE
  }
}
```

```
async fn delete_book(Extension(cnn): Extension<SqlitePool>, Path(id): Path<i32>) -> StatusCode {
  if crate::db::delete_book( connection_pool: &cnn, id).await.is_ok() {
    StatusCode::OK
  } else {
    StatusCode::SERVICE_UNAVAILABLE
  }
}
```



Unit Test the REST System: Framework

Start by building a test framework, and a `setup_tests` function that reads the connection URL, creates a database pool.

Using the `TestClient` from Axum, it creates a local test-version of the website and returns the client - ready for testing.

```
#[cfg(test)]
mod test {
    use super::*;
    use axum_test_helper::TestClient;

    5 usages  ▸ Herbert Wolverson
    async fn setup_tests() -> TestClient {
        dotenv::dotenv().ok();
        let connection_pool : SqlitePool = crate::init_db().await.unwrap();
        let app : Router = crate::router(connection_pool);
        TestClient::new(svc: app)
    }
}
```





Unit Testing the REST Framework

Using the test client, and `tokio::test` for an easy async testing environment - you can quickly test that each of your REST verbs works.

Let's take a quick look at the code for the other tests.

Good news: All the unit tests work!

```
#[tokio::test]
async fn get_all_books() {
    let client : TestClient = setup_tests().await;
    let res : TestResponse = client.get( url: "/books").send().await;
    assert_eq!(res.status(), StatusCode::OK);
    let books: Vec<Book> = res.json().await;
    assert!(!books.is_empty());
}
```





Part 4: Web View (JS+HTML)





Serve static content

Add a new file, `view.rs` and `mod view;` to your main file.

The view system embeds HTML at compile time, and serves it. You could easily use Tower's `ServeDir` feature for dynamic files.



```
use axum::response::Html;
use axum::routing::get;
use axum::Router;

1 usage  👤 Herbert Wolverson
pub fn view_service() -> Router {
    Router::new().route( path: "/", method_router: get( handler: index_page))
}

1 usage  👤 Herbert Wolverson
const INDEX_PAGE: &str = include_str!("index.html");

1 usage  👤 Herbert Wolverson
async fn index_page() -> Html<&'static str> {
    Html(INDEX_PAGE)
}
```

Try it out!

Run the server with cargo run

Navigate to <http://localhost:3001/>

You can add/edit/remove books.



All Books

| # | Author | Title |
|---|--------------------|--------------------|
| 1 | Wolverson, Herbert | Hands-on Rust |
| 2 | Wolverson, Herbert | Rust Brain Teasers |

Book Details

Author

Title

Add Book

Author



Part 5: Containerize in Docker



Build a docker file

Run “docker init”

Select Rust

Select server port 3001

Modify your Dockerfile:

- Add the DATABASE_URL environment variable
- Add migrations to the mounts list



```
# Set the DB URL
ENV DATABASE_URL="sqlite::memory:"
```

```
RUN --mount=type=bind,source=src,target=src \
    --mount=type=bind,source=Cargo.toml,target=Cargo.toml \
    --mount=type=bind,source=Cargo.lock,target=Cargo.lock \
    --mount=type=bind,source=migrations,target=migrations \
    --mount=type=cache,target=/app/target/ \
    --mount=type=cache,target=/usr/local/cargo/registry/ \
    <<EOF
```

Build your Container

`docker build -t axum-test .`

Start the image in Docker Desktop (don't forget to forward port 3001!) - navigate to

<http://localhost:3001/>

And up comes the UI!



All Books

| # | Author | Title |
|---|--------------------|--------------------|
| 1 | Wolverson, Herbert | Hands-on Rust |
| 2 | Wolverson, Herbert | Rust Brain Teasers |

Book Details

Author

Wolverson, Herber

Title

Hands-on Rust

Save

Delete

Add Book

Author

New Author

New Title

Add Book

So far, so good!

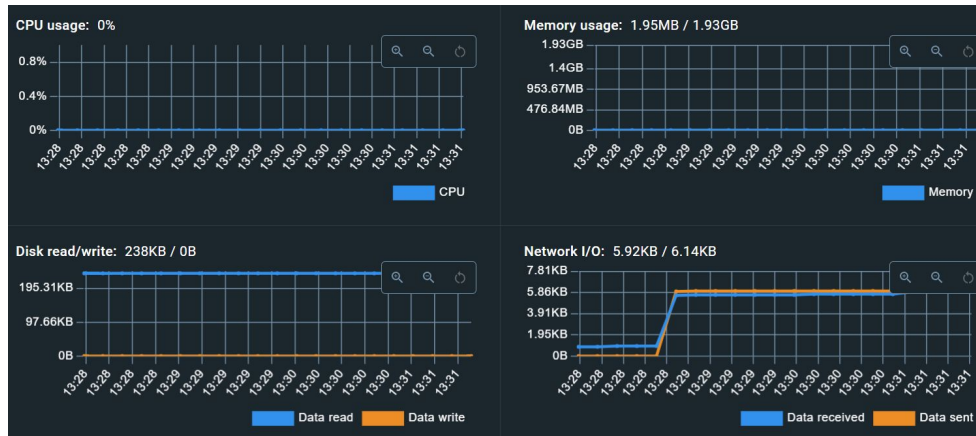
In relatively short time, we've:

- Constructed a database layer with migration support
- Built a database model with a full range of CRUD operations and unit tests
- Created a REST API and fully unit tested it
- Created an HTML/JS view of the API in action

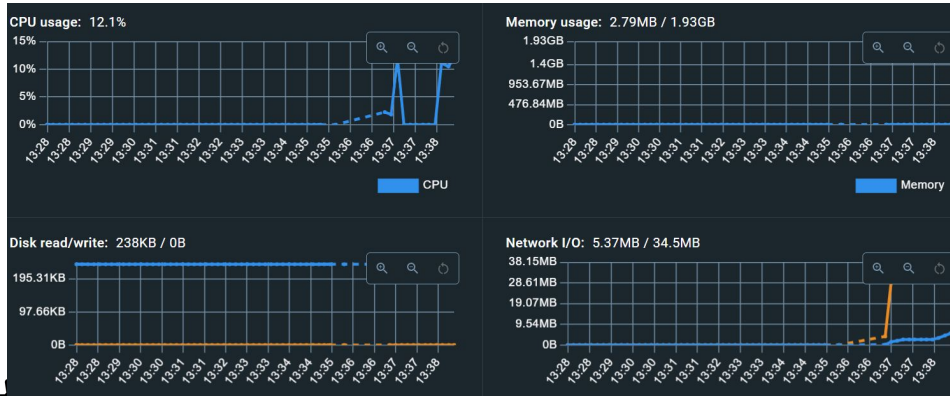
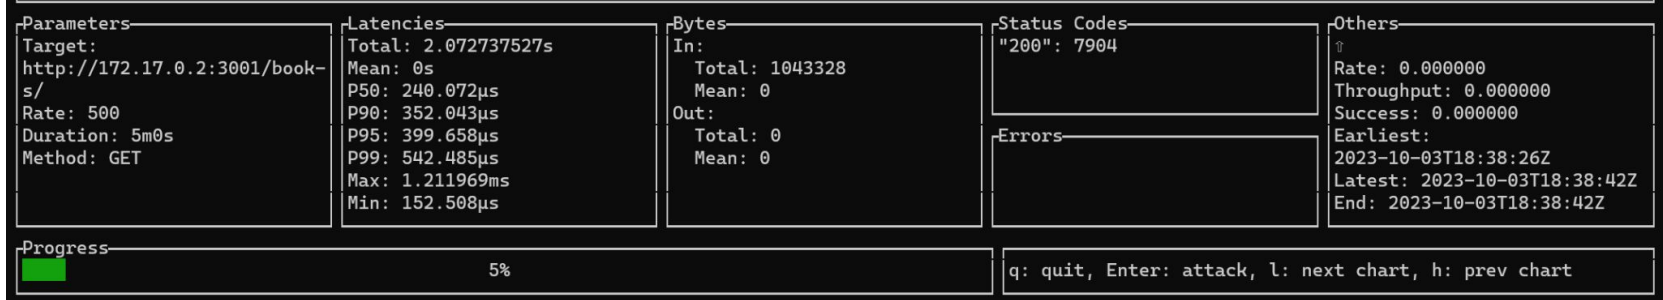


So how resource intensive is this quick application?

While mostly idle:



500 Requests per Second



We're hardly using any memory (3 Mb!)

CPU is ticking upwards - and its mostly SQLite

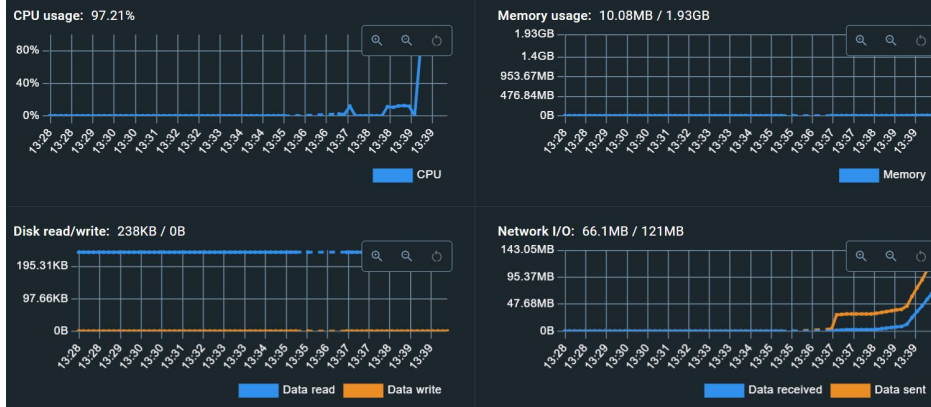
My little MacBook Air can handle this. :-)



At 5,000 Requests per Second



| Parameters | Latencies | Bytes | Status Codes | Others |
|---|---|--|-----------------------------|--|
| Target: http://172.17.0.2:3001/books/ Rate: 5000 Duration: 5m0s Method: GET | Total: 58.801996113s Mean: 0s P50: 192.188µs P90: 315.016µs P95: 539.06µs P99: 2.315229ms Max: 9.989046ms Min: 104.259µs | In: Total: 28134876 Mean: 0 Out: Total: 0 Mean: 0 | "200": 213143 Errors | Requests: 213143 Rate: 0.000000 Throughput: 0.000000 Success: 0.000000 Earliest: 2023-10-03T18:39:29Z Latest: 2023-10-03T18:40:12Z End: 2023-10-03T18:40:12Z |



We're up to 10Mb of RAM! Oh no, however will we host that?

My Macbook Air is struggling a little at 97% CPU usage - but we haven't generated any errors yet.





Let's Spend Some Cache





Let's add a Cache

Hitting the database for a full-table retrieval for every request isn't really necessary.

Rust is supposed to be all about Fearless Concurrency, so let's make a self-invalidating, thread-safe cache for our database model layer.

The basic cache is pretty simple: a Read-Write lock wrapping a list of books, with invalidate and refresh methods.



```
struct BookCache {
    all_books: RwLock<Option<Vec<Book>>>,
}

new *
impl BookCache {
    new *
    fn new() -> Self {
        Self {
            all_books: RwLock::new(value: None),
        }
    }

    new *
    async fn all_books(&self) -> Option<Vec<Book>> {
        let lock : RwLockReadGuard<Option<Vec<Book>>> = self.all_books.read().await;
        lock.clone()
    }

    new *
    async fn refresh(&self, books: Vec<Book>) {
        let mut lock : RwLockWriteGuard<Option<Vec<Book>>> = self.all_books.write().await;
        *lock = Some(books);
    }

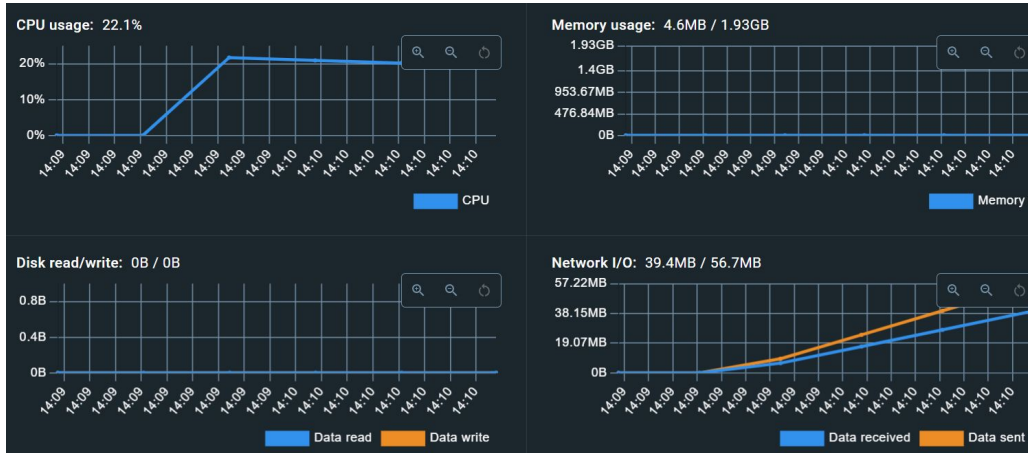
    new *
    async fn invalidate(&self) {
        let mut lock : RwLockWriteGuard<Option<Vec<Book>>> = self.all_books.write().await;
        *lock = None;
    }
}

5 usages new *
static CACHE: Lazy<BookCache> = Lazy::new(|| BookCache::new);
```

Now how does it perform?



| Parameters | Latencies | Bytes | Status Codes | Others |
|---|---|--|-----------------------------|---|
| Target: http://172.17.0.2:3001/books/ Rate: 5000 Duration: 5m0s Method: GET | Total: 17.266785326s Mean: 0s P50: 116.699µs P90: 179.094µs P95: 244.158µs P99: 960.322µs Max: 5.144227ms Min: 26.28µs | In: Total: 15389352 Mean: 0 Out: Total: 0 Mean: 0 | "200": 116586 Errors | ↑ Requests: 116586 Rate: 0.000000 Throughput: 0.000000 Success: 0.000000 Earliest: 2023-10-03T19:09:47Z Latest: 2023-10-03T19:10:11Z End: 2023-10-03T19:10:11Z |



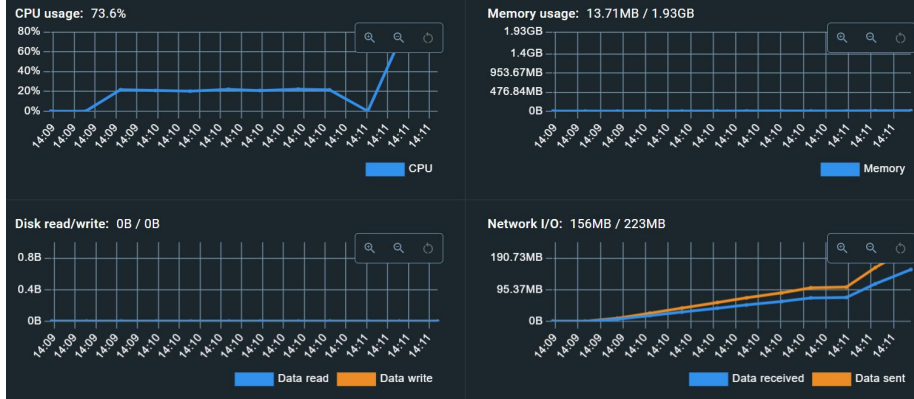
Now the Macbook Air is handling 5,000 requests per second without breaking a sweat - and still only using 5Mb of RAM!



How about 20,000 requests per second now?



| Parameters | Latencies | Bytes | Status Codes | Others |
|--|---|---|-----------------------------|--|
| Target: http://172.17.0.2:3001/books/ Rate: 20000 Duration: 5m0s Method: GET | Total: 17m49.93825839s Mean: 0s P50: 94.652µs P90: 1.154262ms P95: 5.396124ms P99: 28.881083ms Max: 281.897081ms Min: 22.074µs | In: Total: 112370412 Mean: 0 Out: Total: 0 Mean: 0 | "200": 851291 Errors | Requests: 851291 Rate: 0.000000 Throughput: 0.000000 Success: 0.000000 Earliest: 2023-10-03T19:11:04Z Latest: 2023-10-03T19:11:46Z End: 2023-10-03T19:11:46Z |



We're hitting the upper limits of what we can run on a Macbook Air! Latencies are starting to worsen, CPU usage is high. But no errors!

In other words, we can survive bursts - but would need to scale up from a tiny laptop to handle this much live load.



What We've Learned

- SQLX provides a small, fast and expressive database layer with and migrations, test fixtures.
- Axum provides an expressive, enterprise-ready foundation for web services.
- Rust integrates easily with Docker for containerized builds - ready for your infrastructure.
- Rust glues it together nicely, and provides easy concurrency for things like caches.
- Rust is a great choice for line-of-business REST apps.





Next Steps

There's always improvements to be made! Here are some that you could make with relatively little effort - but not in this webinar.

Most of the content covered today is included in *Ultimate Rust: Foundations*.

- Tracing for logging and span timing.
- Tracing for OpenTelemetry support (1 crate and 1 line of code!)
- A proper database!
 - SQLX supports MySQL, PostgreSQL, MS-SQL.
 - Rust drivers are available for most databases.
- A proper cache layer.

